

Threads

Objectives

- Understand the concept of multithreading
- Understand the life cycle of a thread
- Understand thread synchronization
- Be familiar with thread scheduling and prioritization techniques
- Get acquainted with important utility classes for concurrency in the JDK

Threads

- Provide
 - Concurrent paths of execution
 - A way to enable of parallel processing
 - Multi processor/core offer true parallelism
 - Single processor/core give the illusion
- Why Use Threads?
 - Improve program performance
 - Execute separate but coordinated tasks
 - Run background tasks

Java Programs are Multi-Threaded

- Every program uses multiple threads
 - The “main” thread executes `main()`
 - Another thread performs garbage collection
- Many library classes use threading techniques to improve perceived performance and protect data integrity

Constructs

- Two important language constructs support threading:
 - `java.lang.Thread`: Class that represents a thread of execution
 - `Java.lang.Runnable`: Interface that indicates its implementor can do something when called by a thread

java.lang.Thread

- `start()` – housekeeping, becomes runnable
- `run()` – where your code is executed
- `interrupt()` – sets interruption status
 - If blocked on `wait`, `sleep` or `join` `InterruptedException` is thrown and the status cleared
- `isInterrupted()` – returns interruption status
- `interrupted()` – returns interruption status of current thread and resets the status
- `join()` – waits till the thread terminates
- `sleep()` – sleeps for some time
- `yield()` – pause and let another execute
- `set(Default)UncaughtExceptionHandler()` – installs a handler for uncaught exception

java.lang.Thread

- `resume()` - resumes a suspended thread
- `stop()` - stops the thread
- `suspend()` - suspends the thread

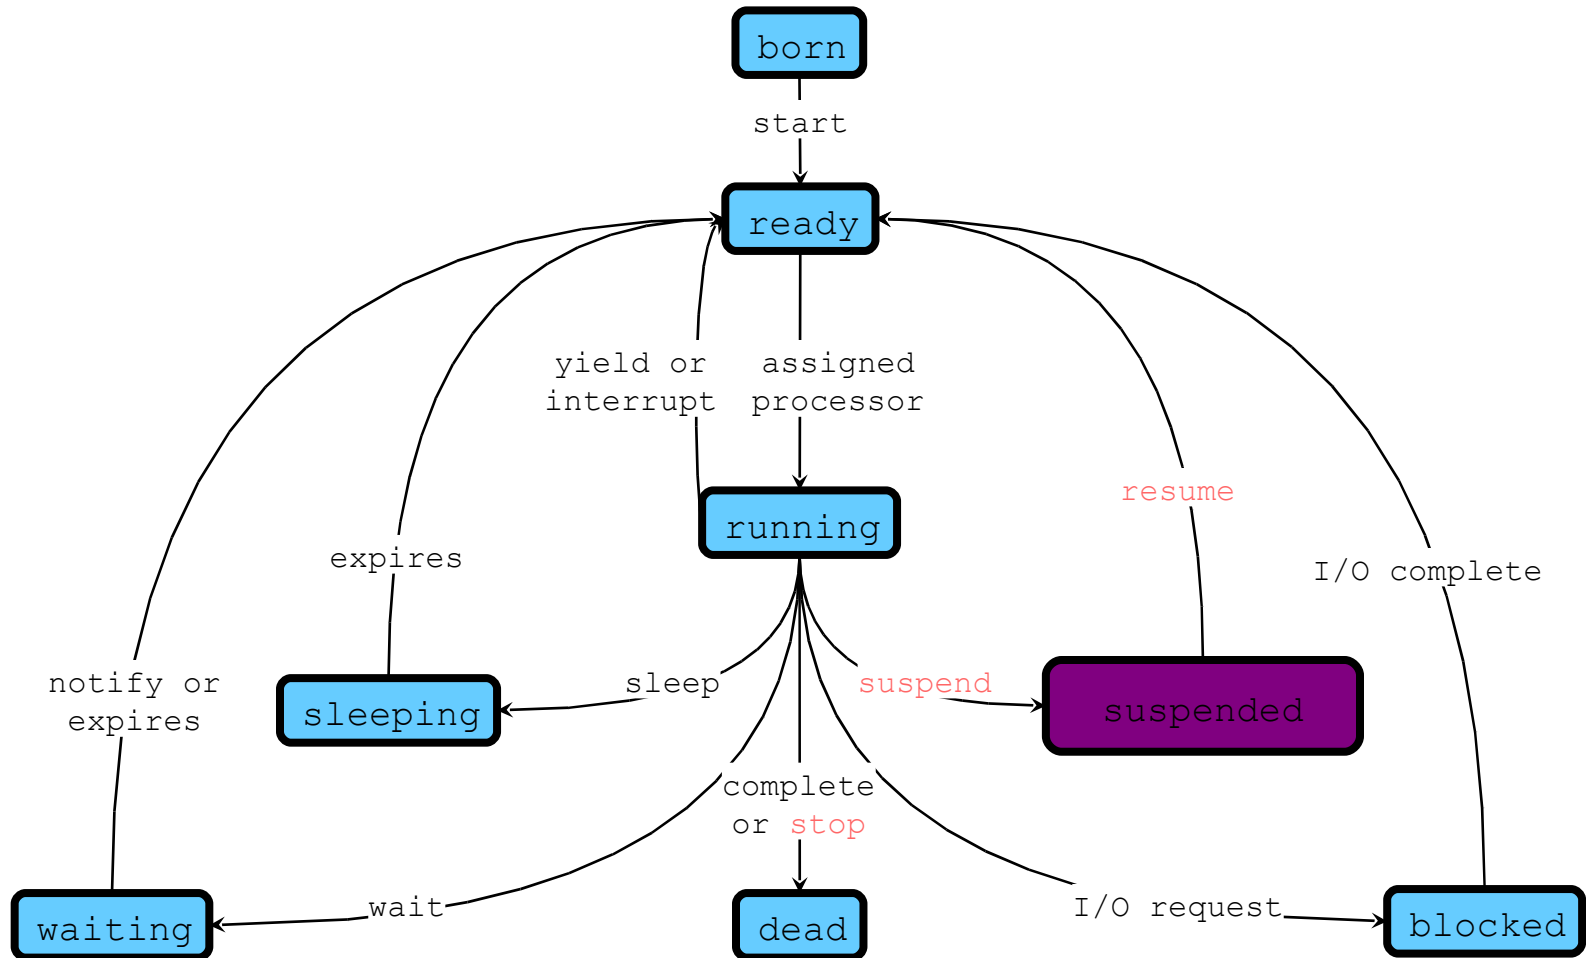
How Threads Run

- Each thread needs time to execute
 - Pre-emptive time-slicing
 - Performed at process level
 - Scheduling
 - No scheduling guarantees
 - No fairness guarantees
 - No guarantee threads will make progress

Thread Priority

- Only thread scheduler uses priority
 - Priority heuristically influences scheduler
 - 10 Crisis management
 - 7-9 Interactive, event driven
 - 4-6 IO - bound
 - 2-3 Background computation
 - 1 Run only if nothing else can
- `setPriority()`
 - `MIN_PRIORITY, 1`
 - `MAX_PRIORITY, 10`
 - `NORM_PRIORITY, 5`

Thread Lifecycle



Controlling Thread State

- Use flag to indicate if thread is running
- Don't be selfish
- Do not use
 - `stop()`
 - `suspend()`
 - `resume()`
- Implement stop using flags, `interrupt()`, `isInterrupted()` and `interrupted()`
- Name threads (useful for debugging)

Controlling Thread State

```
volatile boolean running = true;

public void run()
{
    while( running && !Thread.interrupted() )
    {
        // run and do something in a loop
    }
}

public void stopThisThread()
{
    running = false;
    interrupt();
}
```

Non-portable Thread Features

- The use of priorities and `yield()` to “tune” concurrency should be avoided
- Behavior of `yield()` and impact of priorities vary from one VM to another
- If you have to use `yield()` for performance, it may indicate a higher level design flaw

`java.lang.Runnable`

- Interface specifying the `run()` operation
 - Implemented by `Thread`
- Using `Runnable`
 - Do not need to extend from **`Thread`**
 - Separates logic from concurrency
- Implementation
 - Define implementing class
 - Provide an instance to the `Thread` constructor
 - Start the thread

Synchronization

Reasons for Synchronization

- Example

```
int totalAmount = price * shares;  
// 1. report total amount to the user  
// 2. get money equal to total amount  
// 3. buy shares
```

*Price can be changed by other threads in
between operations
Synchronize access to price*

Monitor and Synchronization

- Monitor (one per object)
 - Implemented in `Object`
 - Only one thread may have ownership of objects monitor at a time
- Synchronization
 - An objects monitor is used to restrict access to guarded statements
 - `synchronized` keyword
 - Method modifier
 - Controls access to a block

synchronized

- A thread may not execute synchronized statements until it obtains ownership of controlling monitor
 - Class methods
 - `Class` object for the method's class
 - Instance methods
 - The object the method was invoked on
 - Synchronized block
 - The object specified in the synchronized block

synchronized Block

```
Object m_Lock = new Object();
public void doSomething()
{
    synchronized(m_Lock)
    {
        // current thread now has ownership of m_Lock
    }
    // current thread releases ownership of m_Lock
    ...
}
```

Atomic Assignment

- Assignment cannot be interrupted

```
foo = 50;  
foo = 42;
```

 - Any thread will see either 42 or 50, not garbage
 - No guarantee value is result of most recent assignment
 - `volatile` – value is “flushed” immediately after written
- Except... **long** and **double**
 - Bigger (64 bit) than native hardware reads/writes
 - Unless `volatile`

Assignment & Synchronization

```
private Object mFoo;  
...  
public void set(Object foo)  
{  
    mFoo = foo;  
}
```

- Doesn't need to be synchronized:
 - Assignment of a reference is atomic
 - ...however, different threads can change the value of mFoo frequently—is this what you want?

Assignment & Synchronization

```
private Object mFoo;  
private boolean mFooAssigned;  
...  
public synchronized void set(Object foo)  
{  
    mFoo = foo;  
    mFooAssigned = true;  
    Syetem.out.println( foo );  
}
```

- Needs to be synchronized, because:
 - Multiple related variables
 - Multiple operations

Assignment & Synchronization

```
public void set(Object foo)
{
    mFoo = foo;

    Iterator it = colListeners.iterator();
    while(it.hasNext())
    {
        ((IEventListener)it.next()).processEvent(mFoo);
    }
}
```

- Do you see a problem here?
 - Notification of listeners can be late
 - Member mFoo can be changed while updating listeners
 - Soapbox (again) be wary of setters, especially in the presence of multiple threads

Assignment & Synchronization

```
public void set(Object foo)
{
    mFoo = foo;

    Iterator it = colListeners.iterator();
    while(it.hasNext())
    {
        ((IEventListener)it.next()).processEvent(foo);
    }
}
```

- Local variables can resolve some problems
 - But spotting such subtle situations is difficult

Cooperating Threads

Monitor Methods

- `java.lang.Object`
 - `wait()`
 - Thread must have ownership of object's monitor prior to calling
 - Thread releases ownership, and waits, until...
 - `notify()`
 - Thread must own object's monitor
 - Thread wakes up **any one** of the threads that called `wait()` on the object
 - Awakened thread must obtain ownership of the object's monitor

Synchronization Example


```
int m_value = 0;
Object m_lock = new Object();
void method1()
{
    synchronized(m_lock)
    {
        ++m_value;
        try {
            m_lock.wait();
        }
        catch(InterruptedException ie) {
            // means something called interrupt()
        }
        finally {
            --m_value;
        }
    }
}
```

Synchronization Example

```
void method2 ()
{
    synchronized(m_Lock)
    {
        m_Lock.notify();
        --m_value;
    }
}
```

What is a value of `m_value` going to be?

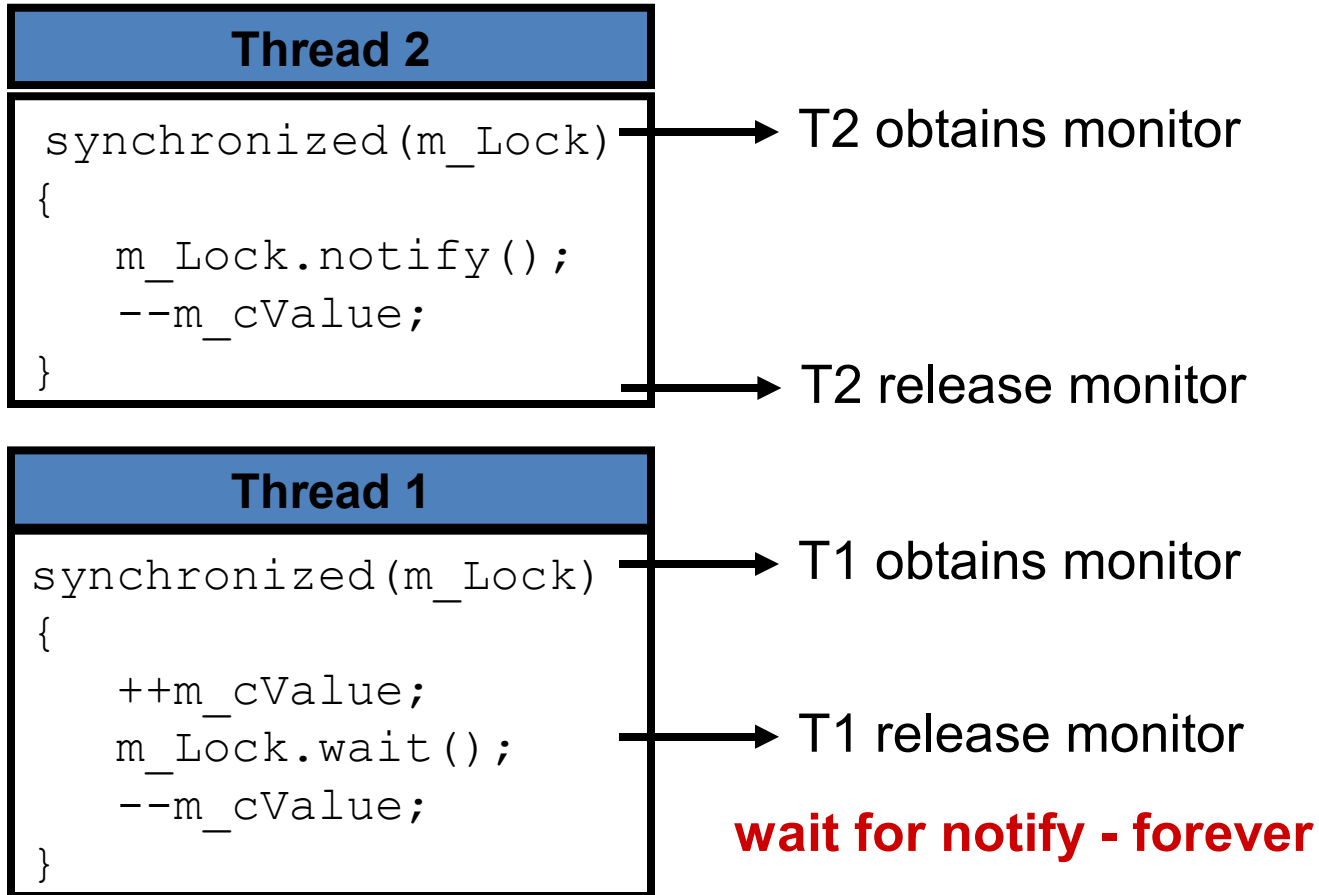
How it Works



More Monitor Methods

- `java.lang.Object`
 - `notifyAll()` - awakens all *wait()*ing threads
 - Easy but more expensive
 - `wait(long timeout)`
 - Equivalent of **`wait()`**, but will try to re-obtain the lock and resume execution
 - `wait(long timeout, int nanos)`
 - Same as above with “higher precision”

Deadlock



Deadlock Tips

- Call `wait()` carefully, without appropriate `notify()` we have deadlock
- Calling `notifyAll()` can help avoid deadlock
- Only good design will save you! :-)

Coordinating Worker Threads

Producer - Consumer

- Both threads work on a specific task
 - Producer thread populates queue
 - Consumer picks the results

```
// consumer thread
public void run()
{
    while( true )
    {
        synchronized(list)
        {
            while(list.isEmpty())
                list.wait()
            Object task = list.getFirst();
        }
    }
}
```

Coordinating Worker Threads

Producer - Consumer

- Producer thread
 - add
 - notify

```
// producer thread
public void run()
{
    synchronized(list)
    {
        list.addLast(task);
        list.notify();
    }
}
```

Coordinating Worker Threads

AND-style

- Workers need to coordinate their work
 - Both tasks must complete prior to proceeding
 - AND-style joiner

```
Thread th1 = new Thread(runnable1);
Thread th2 = new Thread(runnable2);

th1.start();
th2.start();
// both threads running

// wait for first thread to terminate
th1.join();

// wait for second thread to terminate
th2.join();
```

Coordinating Worker Threads

OR-style

- Using callback
 - Each thread knows about others
 - When thread completes notifies other threads
 - Notified threads terminate
 - All joins cease blocking
- Use same `join()` scheme

JDK Concurrency Support

- There are lots of interesting and useful classes in the JDK
 - Synchronized collections (wrapped)
 - “Concurrent” collections
 - BlockingQueue
 - Executor framework

Synchronized Collections

- You don't have to synchronize access to a collection yourself
- The Collections class has static methods that wrap existing collections with synchronization
- Pros:
 - Easy to synchronize access to a collection
- Cons:
 - Large collection + lots of threads = performance degradation

Concurrent Collections

- `java.util.concurrent` has a lot of goodies for concurrency
- Classes with “Concurrent” in the name, e.g. `ConcurrentHashMap` are designed and implemented to support highly concurrent writes and reads:
 - Not possible to lock the entire collection
 - Uses “striped” locking, with a default value of 16 (but you can give any value you like)

BlockingQueue for Concurrency

- Traditional queue that supports four different semantics:
 - Throw an exception
 - Return a special value (boolean)
 - Block until the condition is true
 - Time out after some specified time

	Throw exception	Return value	Block	Time Out
Insert	add(e)	offer()	put(e)	offer(e, time, unit)
Remove	remove()	poll()	take()	poll(time, unit)
Examine	element()	Peek()	n/a	n/a

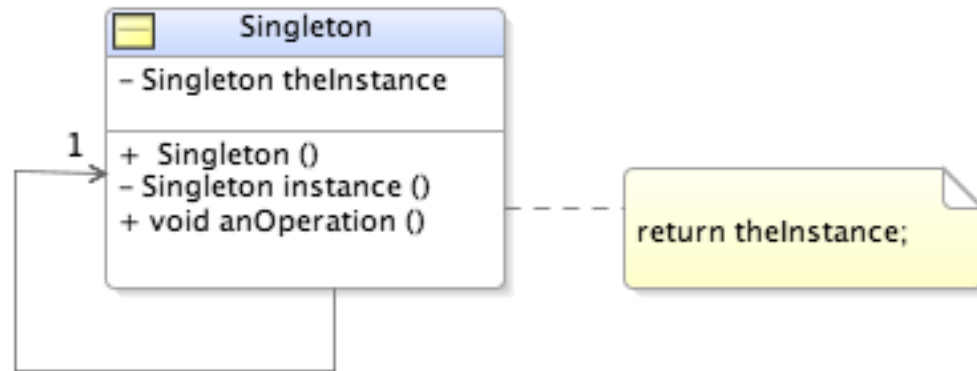
Executor Framework

- Executor is an interface
- Decouples task creation from execution
 - Allows easy abstraction of concurrency algorithms used

Advice on Threading

- It is harder than you think
- If you think you have to use concurrency, start with the prebuilt classes (`java.util.concurrent`)—what you need is probably already there
- Writing your own low level Thread code is a last resort (you are asking for trouble)
- Coding is half the battle; when writing concurrent code there is often much testing, debugging, and tuning to be done

Singleton Pattern



- Insures there is one instance of the class
 - Constructor is private

Singleton enum Implementation

- Use of an enum is the preferred way to implement singletons

```
public enum EnumSingleton {  
    INSTANCE;  
  
    public void anyMethod() {  
        // implementation  
    }  
}
```

Singleton vs. Class Methods

- Benefits over a class having only class (static) methods
 - Constructors apply to instances
 - Only instances can be serialized
 - Interfaces can't specify static methods
 - Inheritance is not useful
 - May insert an object pool if a single instance is not sufficient (non-enum implementation)

Lazy Initialization Class Holder Idiom

- Lazily initializes a static field
 - Classes are guaranteed not to be initialized until used
 - Useful for singleton implementation pre 1.5

```
public class MySingleton {  
    private static class InstanceHolder {  
        static final MySingleton singleton = new MySingleton();  
    }  
  
    private MySingleton() {  
    }  
  
    public static MySingleton instance() {  
        return InstanceHolder.singleton;  
    }  
}
```

Lazy Initialization

Double-check Idiom

- Double-check idiom
 - Broke prior to 1.5, due to weak `volatile` semantics

```
public class LazyInitialization {
    private volatile LazyType lazyValue;

    public LazyType getLazyValue() {
        LazyType result = lazyValue;
        if (result == null) {
            synchronized(this) {
                if (result == null) {
                    result = lazyValue = computeInitValue();
                }
            }
        }
        return result;
    }
    ...
}
```

Misc

Shutting Down a Program

- To clean up and shut down a running Java program from the outside (from the command line or a script file)
- Use a *shutdownHook*
 - Simply an initialized but unstarted thread
 - *Must* extend Thread, Runnable doesn't work here
 - All shutdownHooks will be run when the Java runtime is told to stop; by a Cntl-C from the command line, for example

Shutting Down a Program (cont'd)

- Instantiate the shutdownHook
- Get the current runtime object by calling `Runtime.getRuntime();`
- Register the shutdownHook by calling `runtime.addShutdownHook(shutdownHook);`
- When the runtime shuts down, all the registered shutdownHooks will run